

pHExploration


Alexandra Lawrence

In order to understand the variable that I want to explore, I want to look at the distribution of pH along with any patterns it may follow over time.

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

```
## Warning: Removed 48 rows containing non-finite values (stat_bin).
```


Distribution of pHMEAN



Min	Max	Mean	Median	Q1	Q3	Standard_Deviation	IQR
7.448	8.257	8.002	7.995	7.935	8.078	0.098	0.143

According to the histogram, pHMEAN appears to be bimodal and very slightly left-skewed. The average pH appears to be slightly basic, with a value of 8.002. The range of pH between the first and third quarters is 0.14. Additionally, on average, every value has about a 0.098 average distance from the mean.

```
## Warning: Removed 48 rows containing missing values (geom_point).
```



It is clear that the pH levels off of the coast vary throughout the year, with a higher pH in the colder months, and a lower pH value in the warmer months. Additionally, it looks as though the lowest value gets lower every year, but this will have to be examined in further detail. Values recorded during a named storm also appear to be on the smaller end, so correlations will be examined in greater detail.

```
##      pHMEAN      Date
## 1 7.448091 2018/09/20
```

The smallest pH value was recorded on 9/20/2018, which was about a week after Hurricane Florence hit North Carolina. Therefore, we will look into the influence of storms as well as how long these effects might last. However, is this because of the storm or just a coincidence because pH tends to lower in warmer seasons and storms happen more often in summer?

```
## Warning: Removed 48 rows containing non-finite values (stat_boxplot).
```

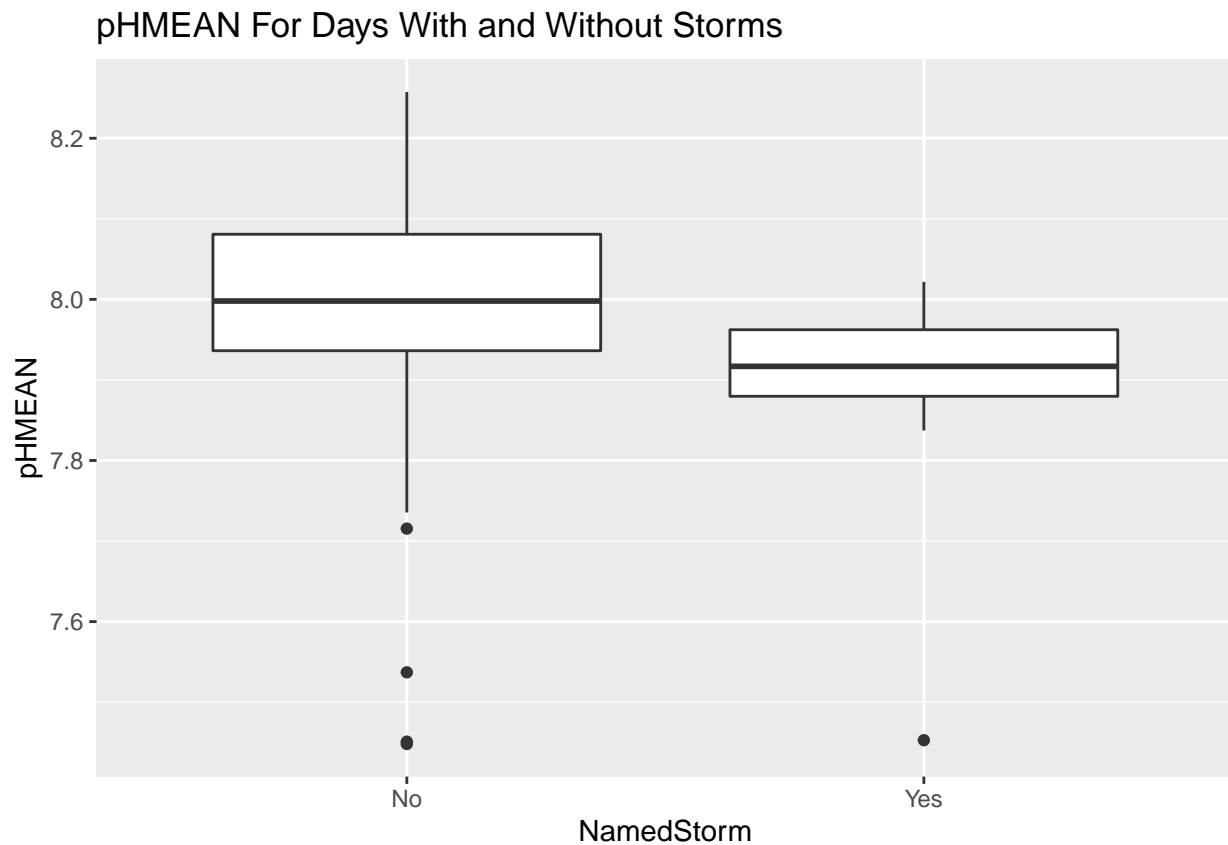




Table 2: Pairwise Wilcox Test

group1	group2	p.value
Yes	No	9e-05

According to the boxplot and pairwise wilcox test, with a p-value of 0.00009, there is in fact a difference in the mean pH value between values recorded during a named storm and those recorded during typical weather.



```
## Warning: Removed 48 rows containing missing values (geom_point).
```



There appears to be a somewhat negative linear relationship between temperature and pH, according to this plot. As the mean temperature increases, the mean pH decreases. This makes sense in the context of the data. Additionally, according to the correlation plot there is a relatively strong negative correlation between these two variables.


Table 3: Linear Model of pH and Temperature

term	estimate	std.error	statistic	p.value
(Intercept)	8.2140314	0.0076408	1075.01666	0
TemperatureMEAN	-0.0104838	0.0003566	-29.40315	0

For every one degree increase in temperature, the pH is estimated to decrease by 0.0105 on average. The p-value is close to zero, meaning that there is a relationship between pH and temperature.


```
## Warning: position_stack requires non-overlapping x intervals
```

```
## Warning: position_stack requires non-overlapping x intervals
```




##	Date	pHMEAN	diff	NamedStorm
## 1	2010/12/08	8.257426	0.2554793	No
## 2	2010/12/15	8.248694	0.2467472	No
## 3	2010/12/22	8.229639	0.2276926	No
## 4	2010/12/29	8.233435	0.2314883	No
## 5	2011/01/05	8.207517	0.2055706	No
## 6	2011/01/12	8.253215	0.2512684	No
## 7	2011/01/19	8.189668	0.1877210	No
## 8	2011/01/26	8.191532	0.1895856	No
## 9	2011/01/27	8.180976	0.1790294	No
## 10	2016/01/20	8.189257	0.1873104	No
## 11	2019/01/16	8.178843	0.1768957	No
##	Date	pHMEAN	diff	NamedStorm
## 1	2012/08/29	7.820919	-0.1810281	No
## 2	2014/08/06	7.715420	-0.2865267	No
## 3	2014/08/20	7.847530	-0.1544166	No
## 4	2015/10/07	7.450993	-0.5509536	No
## 5	2016/10/12	7.846731	-0.1552157	Yes
## 6	2017/09/26	7.847970	-0.1539768	Yes
## 7	2018/08/08	7.825095	-0.1768513	No
## 8	2018/09/19	7.452792	-0.5491547	Yes
## 9	2018/09/20	7.448091	-0.5538559	No
## 10	2018/09/21	7.537025	-0.4649218	No
## 11	2018/09/22	7.735478	-0.2664689	No
## 12	2019/09/07	7.837277	-0.1646693	Yes
## 13	2019/09/10	7.748015	-0.2539317	No


```
## 14 2019/09/11 7.821565 -0.1803822 No
## Warning: position_stack requires non-overlapping x intervals
## Warning: position_stack requires non-overlapping x intervals
```




```
## Warning: position_stack requires non-overlapping x intervals
## Warning: position_stack requires non-overlapping x intervals
```



Additional Variables



Warning: Removed 48 rows containing missing values (geom_point).



There appears to be a negative linear relationship between water-atmospheric carbon dioxide and pH. As the water-atmospheric CO₂ increases, the pH decreases. Additionally, there appear to be more named storms recorded at higher values of water-atmospheric CO₂ increases. According to the correlation plot, there is a strong negative correlation between these two variables.

Table 4: Linear Model of pH and DeltaCO2

term	estimate	std.error	statistic	p.value
(Intercept)	8.0413128	0.0007747	10379.7161	0
deltaCO2	-0.0008537	0.0000064	-133.6373	0

For every 100 μatm increase in water-atmospheric carbon dioxide, the pH is estimated to decrease by 0.0854 on average. The p-value is close to zero, meaning there is a relationship between the two variables.

```
## deltaCO2      Date
## 1 737.8547 2015/10/07
```

The largest value for deltaCO₂ occurs on a day when the second smallest pH is recorded.

```
## Warning: Removed 48 rows containing non-finite values (stat_boxplot).
```

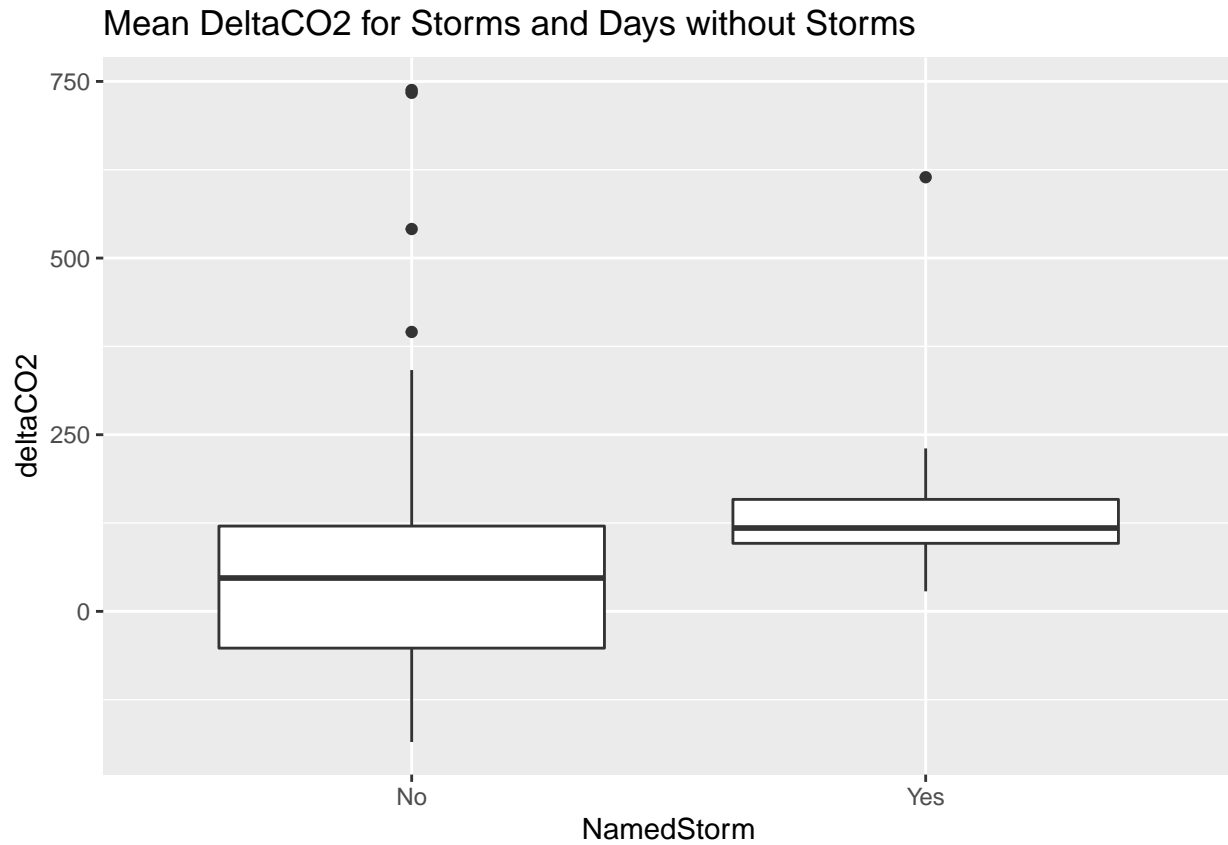





Table 5: Pairwise Wilcox Test

group1	group2	p.value
Yes	No	0.0005083

After examining the boxplot and administering a pairwise wilcox test, with a p-value of 0.0005, it is apparent that there is a difference in the mean water-atmospheric carbon dioxide for days with named storms and days with typical weather.



Warning: Removed 382 rows containing missing values (geom_point).



There does not appear to be a strong relationship between NH4 and pH. The correlation plot also displays a weak correlation between the two variables.

Table 6: Linear Model of pH and NH4

term	estimate	std.error	statistic	p.value
(Intercept)	8.0343996	0.0055237	1454.545014	0
NH4MEAN	-0.0000796	0.0000123	-6.454661	0

For every one nM increase in NH4, the pH decreases by -0.00008. While the small p-value suggests a relationship between the two variables, this may be a result of the influence from the outliers.

Warning: Removed 381 rows containing non-finite values (stat_boxplot).

Mean NH4MEAN For Storms and Days without Storms




Table 7: Pairwise Wilcoxon Test

group1	group2	p.value
Yes	No	0.1976984

According to the boxplot there looks like there is a difference in the mean NH4 for days with a named storm, however, the pairwise wilcoxon test fails to reject the hypothesis that there is no difference in means.

Multiple Linear Regression

In this linear model, I will include the temperature, named storm indicator, water-atmospheric carbon dioxide, and NH4 because each variable appears to have a relationship with pH. I will also include an interaction term between the storm indicator and each of the numeric variables because there is a difference in means for values recorded during a storm.

Table 8: Multiple Linear Regression Model for pH

term	estimate	std.error	statistic	p.value
(Intercept)	8.0414305	0.0034884	2305.1744161	0.0000000
NamedStormYes	-0.0482955	0.0466487	-1.0353009	0.3013939
TemperatureMEAN	0.0002437	0.0001839	1.3253833	0.1860908
deltaCO2	-0.0008351	0.0000128	-65.3884154	0.0000000
NH4MEAN	-0.0000023	0.0000021	-1.0490303	0.2950407
NamedStormYes:TemperatureMEAN	0.0017248	0.0022417	0.7694073	0.4422800
NamedStormYes:deltaCO2	-0.0000229	0.0001669	-0.1370037	0.8911233
NamedStormYes:NH4MEAN	-0.0000070	0.0000104	-0.6755170	0.4998875

Several variables have insignificant p-values in this model, so I will perform backwards selection in order to ensure that I get the most accurate results. This test will eliminate any unnecessary predictors, which is important as any model can get a high r-squared value with a lot of variables, even if many are not considered significant.

Table 9: Reduced Model through backward Selection

term	estimate	std.error	statistic	p.value
(Intercept)	8.0397817	0.0032301	2488.996526	0.0000000
NamedStormYes	-0.0160253	0.0064739	-2.475359	0.0138757
TemperatureMEAN	0.0003153	0.0001755	1.796980	0.0733685
deltaCO2	-0.0008414	0.0000115	-72.866298	0.0000000

The reduced model only includes the named storm indicator, temperature, and water-atmospheric carbon dioxide. While temperature does not have a significant p-value, I will leave it in the model for the moment because of the results from exploratory data analysis. However, this may need to be investigated further.

Table 10: Model Selection Criteria for Full Model


Adjusted R Squared	AIC	BIC
0.9797	-1735.585	-1702.341

Table 11: Model Selection Criteria for Reduced Model

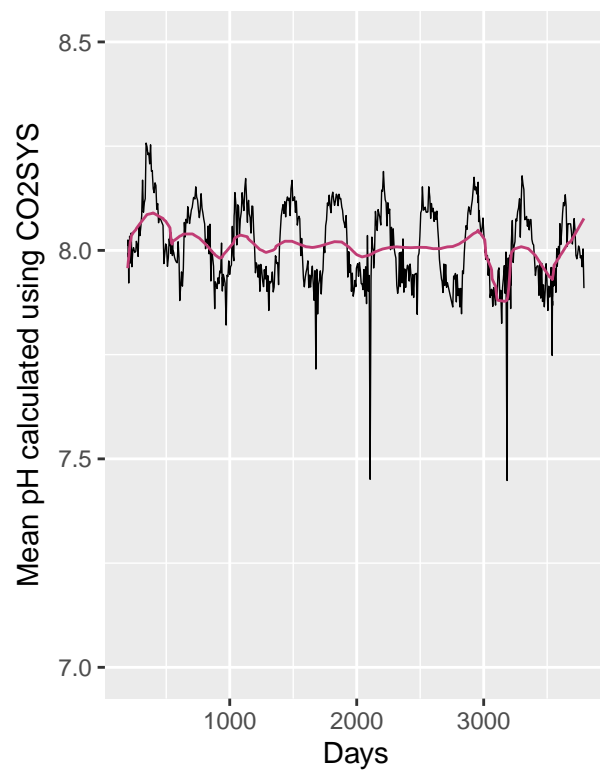
Adjusted R Squared	AIC	BIC
0.9798	-1739.935	-1721.466

The adjusted r-squared, AIC, and BIC values are similar for the full and reduced models. However, the reduced model has a slightly higher adjusted r-squared and lower AIC and BIC values, so I will choose this as the most accurate.

Time Series



Trend Mapping onto Data



Seasonal Cycle Mapping onto Data






Table 12: Seasonal Mann Kendall test for pH

p.value	kendall_score	statistic
0.0021753	-328	-0.0932878





p-value is less than 0.05, so we can reject the null hypothesis – There is a possible trend in the data

Series full_pHMEAN\$pHMEAN

Series full_pHMEAN\$pHMEAN



```
## Series: pico1.ts
## ARIMA(1,1,3)
##
## Coefficients:
##      ar1      ma1      ma2      ma3
##    -0.7071  0.3294 -0.2159 -0.1328
## s.e.   0.1300  0.1343  0.0685  0.0410
##
## sigma^2 estimated as 0.002793: log likelihood=960.33
## AIC=-1910.66  AICc=-1910.57  BIC=-1888.43
```



```
## Series: pico1_seasadj
## ARIMA(0,1,4)(0,0,2)[52]
##
## Coefficients:
##      ma1      ma2      ma3      ma4      sma1      sma2
##    -0.3853  0.0486 -0.1800  0.0887 -0.1576 -0.1442
## s.e.   0.0403  0.0426  0.0397  0.0377  0.0499  0.0579
```

```
##
## sigma^2 estimated as 0.002499: log likelihood=994.37
## AIC=-1974.74 AICc=-1974.56 BIC=-1943.62

Test for Stationarity:

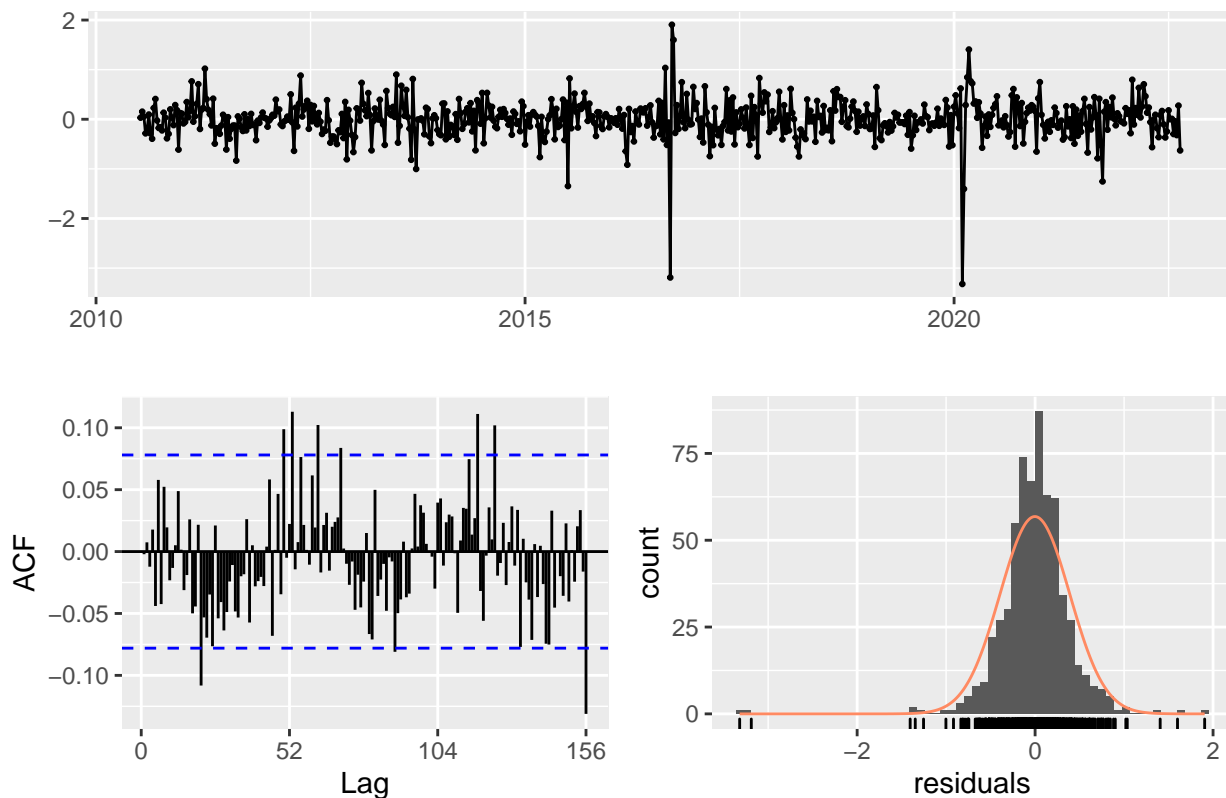
## Warning in kpss.test(pico1_seasadj): p-value smaller than printed p-value

##
## KPSS Test for Level Stationarity
##
## data: pico1_seasadj
## KPSS Level = 0.85542, Truncation lag parameter = 6, p-value = 0.01

P-value is less than 0.05 – could be significant and the data is stationary

## Series: pico1_seasadj
## ARIMA(0,1,4)(0,0,2)[52]
## Box Cox transformation: lambda= 1.999924
##
## Coefficients:
##          ma1      ma2      ma3      ma4      sma1      sma2
##      -0.3857  0.0518 -0.1778  0.0930 -0.1588 -0.1468
## s.e.   0.0403  0.0426  0.0398  0.0377  0.0503  0.0585
##
## sigma^2 estimated as 0.1555: log likelihood=-306.76
## AIC=627.53 AICc=627.71 BIC=658.65
```

Residuals from ARIMA(0,1,4)(0,0,2)[52]



```
##
## Ljung-Box test
```



```
##  
## data: Residuals from ARIMA(0,1,4)(0,0,2)[52]  
## Q* = 7.3346, df = 6, p-value = 0.291  
##  
## Model df: 6. Total lags used: 12  
## [1] 627.7082
```

